

Machine learning 04

Byung Chang Chung

Gyeongsang National University

bcchung@gnu.ac.kr

Contents


- Dimensionality reduction

Curse of dimensionality

- Lots of samples
 - training sample has thousands or even millions of properties
 - not only slow down training but also make it difficult to find a good solution

Curse of dimensionality

- We cannot intuitively imagine high-dimension
 - if you select any two points
 - in 3D, the average distance is approximately 0.66
 - in 1,000,000D, the average distance is approximately 428.25



Curse of dimensionality



- How to approach?
 - theoretically, increasing the size of the training set until the training sample becomes sufficiently dense
 - the number of training samples required to reach a certain density increases exponentially as the number of dimensions increases

Projection

- Definition
 - in mathematics, a projection is a mapping of a set (or other mathematical structure) into a subset (or sub-structure), which is equal to its square for mapping composition

Projection

- Examples




Projection

- Properties of data
 - training samples are not distributed evenly across all dimensions
 - many features remain almost unchanged
 - other properties are strongly associated with each other



Projection

- Properties of data
 - resulting in all training samples lying (or close) in a low-dimensional subspace within a high-dimensional space




Projection

- Properties of data
 - projecting all training samples perpendicular to this subspace yields a below dataset




Projection

- Exceptions
 - Swiss roll dataset




Projection

- Exceptions
 - projected Swiss roll dataset
 - which one is effective projection?



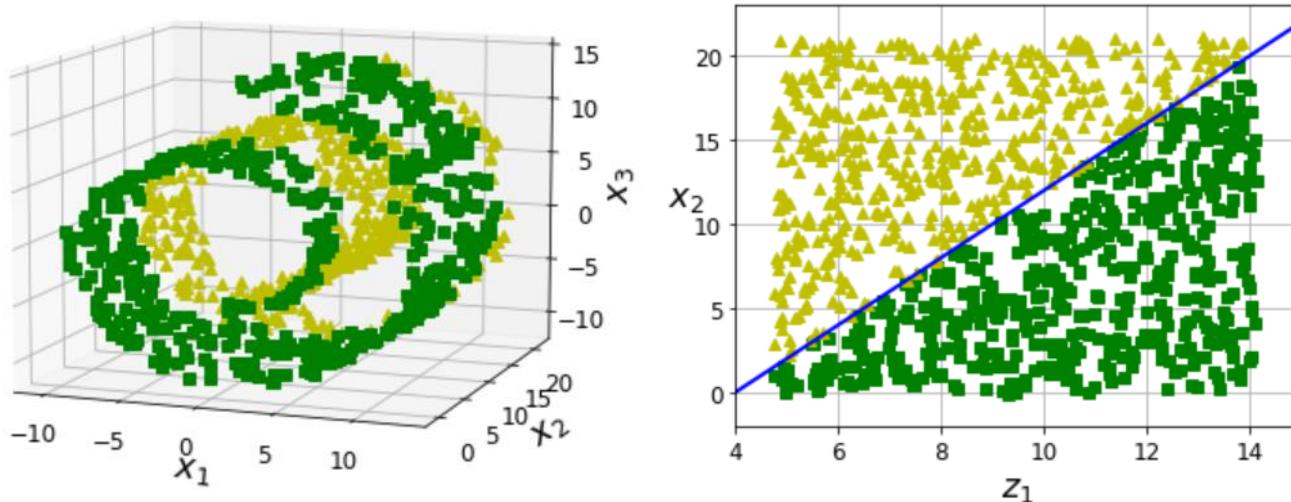
Manifold learning

- Definition of manifold
 - in mathematics, a manifold is a topological space that locally resembles Euclidean space near each point



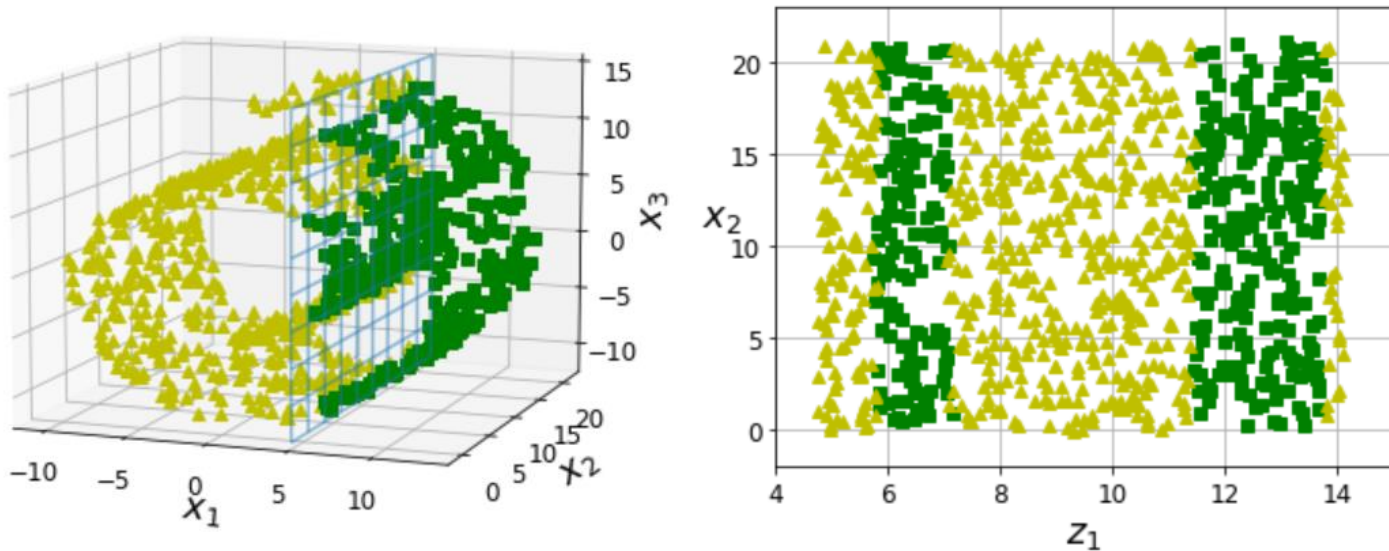
Manifold learning

- Assumption
 - Based on the hypothesis that the high-dimensional dataset lies closer to the low-dimensional manifold




Manifold learning

- Assumption
 - is not always correct




PCA

- Principal component analysis (주성분분석)
 - find hyperplane that closer to the given dataset
 - then project on the optimal hyperplane



PCA process

- How to select optimal hyperplane?
 - in terms of variance
 - dotted line
 - real line maximizes variance



PCA process

- Principal component
 - find axis which results largest variance
 - singular value decomposition in linear algebra

$$A = U\Sigma V^T$$

A : $m \times n$ rectangular matrix

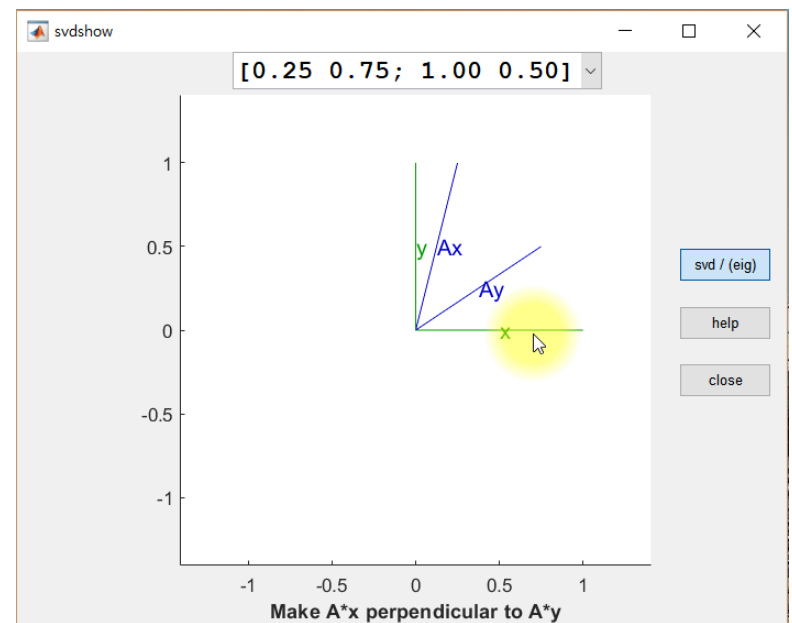

U : $m \times m$ orthogonal matrix

Σ : $m \times n$ diagonal matrix

V : $n \times n$ orthogonal matrix

PCA process

- Singular value decomposition



PCA process

- Projection to d-dimension

$$\mathbf{X}_{d\text{-proj}} = \mathbf{X}\mathbf{W}_d$$

- \mathbf{W}_d refers to the first d-rows of matrix \mathbf{V}

PCA implementation

- Implementation in scikit-learn

```
from sklearn.decomposition import PCA

pca = PCA(n_components = 2)
X2D = pca.fit_transform(X)
```

- reduce dimension to 2 (n_components = 2)

PCA implementation


- Finding appropriate dimension

```
pca = PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
d = np.argmax(cumsum >= 0.95) + 1
```

- finding dimension which maintain the variance of dataset greater than 0.95

PCA implementation

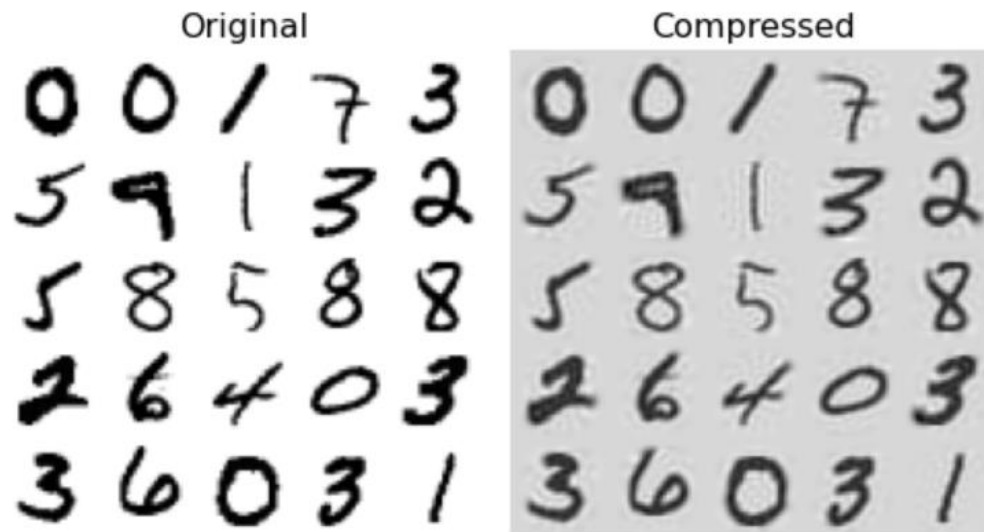
- Finding appropriate dimension – visualization



PCA application

- PCA for extraction
 - reducing dimension => reducing amount of information

$$\mathbf{X}_{\text{recovered}} = \mathbf{X}_{d\text{-proj}} \mathbf{W}_d^T$$




PCA application

- Random PCA
 - finding randomized SVD can find principal component faster
 - SVD complexity $O(m \times n^2) + O(n^3)$
 - randomized SVD complexity $O(m \times d^2) + O(d^3)$

```
rnd_pca = PCA(n_components=154, svd_solver="randomized", random_state=42)  
X_reduced = rnd_pca.fit_transform(X_train)
```

PCA application

- Time complexity of random PCA



PCA application

- Incremental PCA (IPCA)
 - how to make online?
 - divide dataset into minibatch
 - using IPCA

```
from sklearn.decomposition import IncrementalPCA

n_batches = 100
inc_pca = IncrementalPCA(n_components=154)
for X_batch in np.array_split(X_train, n_batches):
    print(".", end="") # 책에는 없음
    inc_pca.partial_fit(X_batch)

X_reduced = inc_pca.transform(X_train)
```

Kernel PCA


- Nonlinear projection using kernel

```
from sklearn.decomposition import KernelPCA  
  
rbf_pca = KernelPCA(n_components=2, kernel="rbf", gamma=0.04)  
X_reduced = rbf_pca.fit_transform(X)
```

- kernel calculation can make non-linearity

Kernel PCA

- Nonlinear projection using kernel – visualization



Kernel PCA

- Finding proper kernel and hyperparameter
 - there are no clear performance metrics for selecting good kernels and hyperparameters since it is unsupervised
 - use grid search to select the appropriate kernel and hyperparameters for a imaginary regression problem

Kernel PCA

- Finding proper kernel and hyperparameter

```
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline

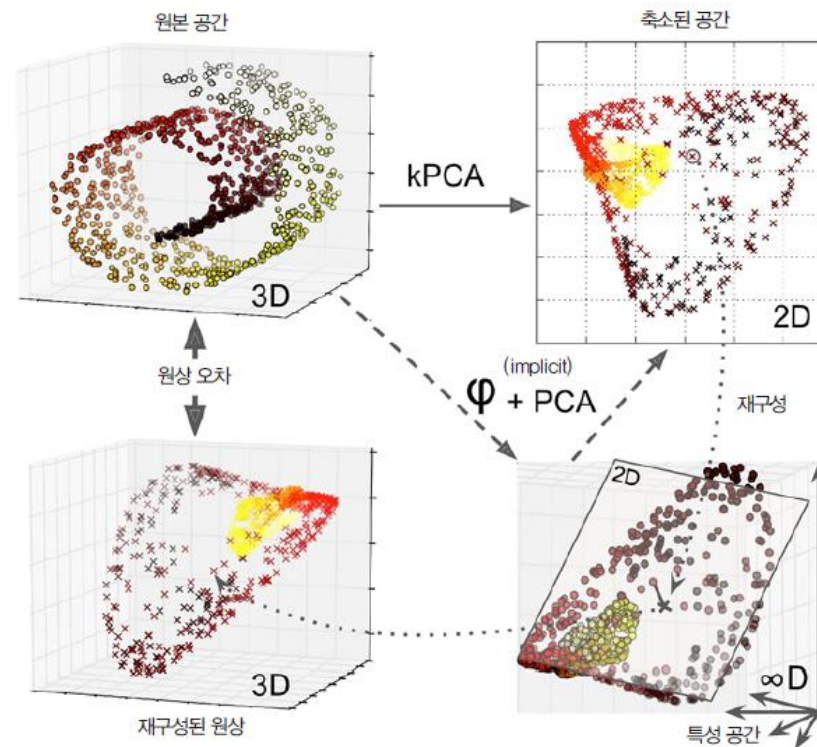
clf = Pipeline([
    ("kpca", KernelPCA(n_components=2)),
    ("log_reg", LogisticRegression(solver="lbfgs"))
])

param_grid = [{
    "kpca__gamma": np.linspace(0.03, 0.05, 10),
    "kpca__kernel": ["rbf", "sigmoid"]
}]

grid_search = GridSearchCV(clf, param_grid, cv=3)
grid_search.fit(X, y)
```

Kernel PCA

- Other way
 - comparison with original data and reconstructed data



Locally linear embedding

- Manifold learning which is not limited to projection
 - measure how linearly each training sample relates to the closest neighbor (c.n.)
 - local relationship can be preserved
 - it works well to unfold the twisted manifold when there is not too much noise

Locally linear embedding

- Mathematical representation


$$\hat{\mathbf{W}} = \underset{\mathbf{W}}{\operatorname{argmin}} \sum_{i=1}^m \left(\mathbf{x}^{(i)} - \sum_{j=1}^m w_{i,j} \mathbf{x}^{(j)} \right)^2$$

$$[\text{조건}] \begin{cases} w_{i,j} = 0 & \mathbf{x}^{(j)} \text{가 } \mathbf{x}^{(i)} \text{의 최근접 이웃 } k \text{개 중 하나가 아닐 때} \\ \sum_{j=1}^m w_{i,j} = 1 & i = 1, 2, \dots, m \text{일 때} \end{cases}$$

$$\mathbf{Z} = \underset{\mathbf{Z}}{\operatorname{argmin}} \sum_{i=1}^m \left(\mathbf{z}^{(i)} - \sum_{j=1}^m \hat{w}_{i,j} \mathbf{z}^{(j)} \right)^2$$

Locally linear embedding

- LLE in scikit-learn



Other works with dimension

- Random projection
 - random linear projections to project data into low-dimensional space
- Multidimensional scaling (MDS)
 - reduce dimension while preserving distance between samples

Other works with dimension


- Isomap
 - make a graph by linking each sample with the nearest neighbor
 - reduce dimension while maintaining the geodesic distance

Other works with dimension

- t-distributed stochastic neighbor embedding (t-SNE)
 - reduce dimensions while keeping similar samples close and non-similar samples far away
- Linear discriminant analysis (LDA)
 - finding hyperplane that best separates the dataset
 - this hyperplane will be used in projection

Other works with dimension

- Results of dimension reduction algorithm



Feel free to question
Through e-mail & LMS

본 자료의 연습문제는 수업의 본교재인
한빛미디어, Hands on Machine Learning(2판)에서 주로 발췌함