

Machine learning 10

Byung Chang Chung

Gyeongsang National University

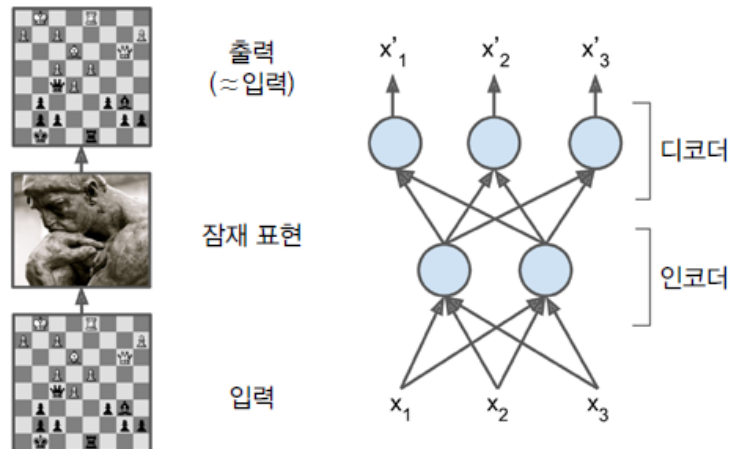
bcchung@gnu.ac.kr

Contents

- Autoencoder and reinforcement learning
 - basic structure of autoencoder
 - types of autoencoder
 - applications of autoencoder
 - basic theory of reinforcement learning
 - applications of reinforcement learning

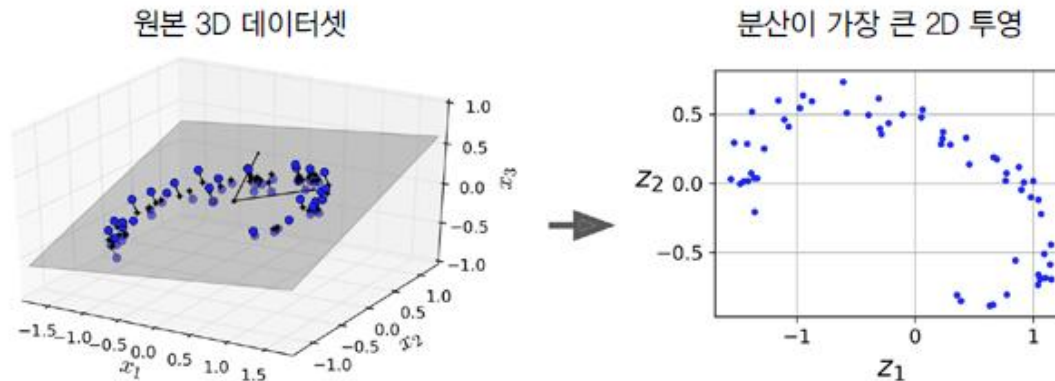
What is autoencoder

- Encoder and decoder
 - the process of encoding converts information from a source into symbols for communication or storage
 - decoding is the reverse process, converting code symbols back into a form that the recipient understands



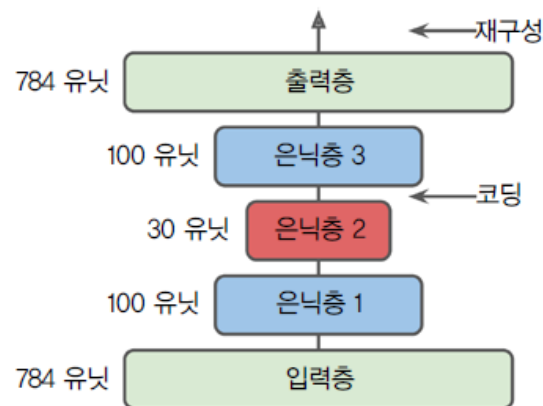
What is autoencoder

- Autoencoder in linear process
 - like PCA
 - encoded results can be viewed similar with the results of dimension reduction



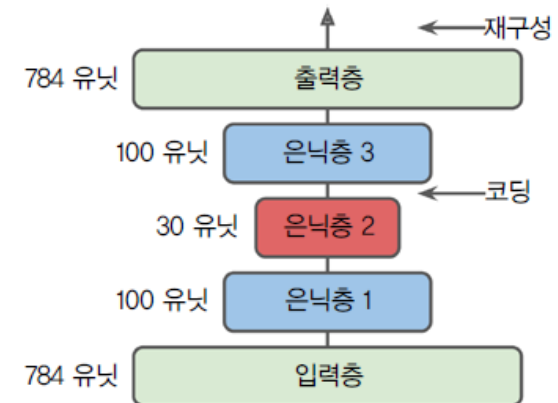
What is autoencoder

- Stacked encoder
 - an autoencoder which has multiple hidden layers
 - compare inputs and outputs to ensure that the autoencoder is properly trained



Applications of autoencoder

- Fashion MNIST dataset
 - usage of stacked autoencoder
 - structure is the same as previous image



```
stacked_encoder = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(100, activation="selu"),
    keras.layers.Dense(30, activation="selu"),
])
stacked_decoder = keras.models.Sequential([
    keras.layers.Dense(100, activation="selu", input_shape=[30]),
    keras.layers.Dense(28 * 28, activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])
stacked_ae = keras.models.Sequential([stacked_encoder, stacked_decoder])
stacked_ae.compile(loss="binary_crossentropy",
                   optimizer=keras.optimizers.SGD(learning_rate=1.5), metrics=[rounded_accuracy])
history = stacked_ae.fit(X_train, X_train, epochs=20,
                        validation_data=(X_valid, X_valid))
```


Applications of autoencoder

- Fashion MNIST dataset
 - by comparing input and output, the training can be validated




Applications of autoencoder

- Fashion MNIST dataset
 - dimension reduction
 - middle hidden layer has 30 dimensions
 - after 30-D, use t-SNE to reduce the dimension to 2



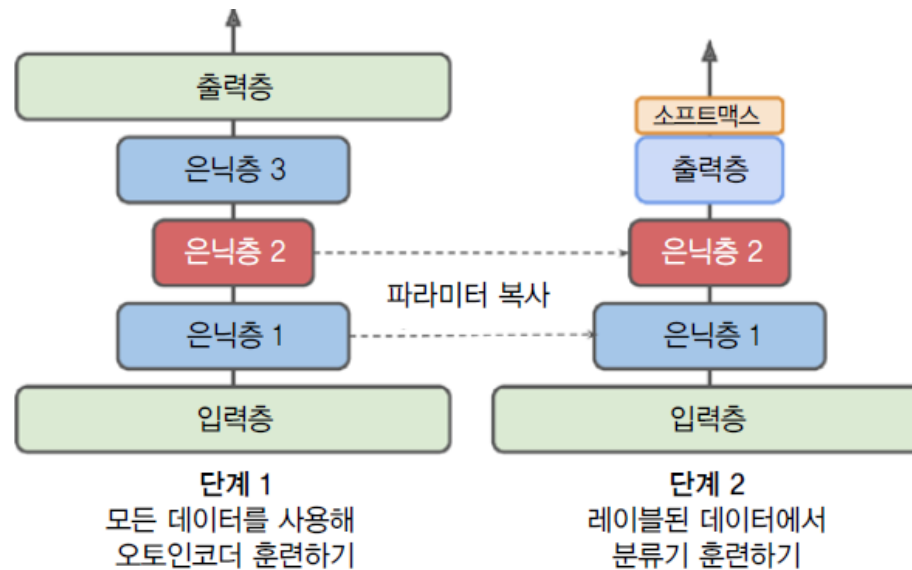
Applications of autoencoder

- Fashion MNIST dataset
 - result of dimension reduction



Applications of autoencoder

- Fashion MNIST dataset
 - unsupervised pre-training



Applications of autoencoder

- Fashion MNIST dataset
 - tied weights

```
class DenseTranspose(keras.layers.Layer):
    def __init__(self, dense, activation=None, **kwargs):
        self.dense = dense
        self.activation = keras.activations.get(activation)
        super().__init__(**kwargs)
    def build(self, batch_input_shape):
        self.biases = self.add_weight(name="bias",
                                      shape=[self.dense.input_shape[-1]],
                                      initializer="zeros")
        super().build(batch_input_shape)
    def call(self, inputs):
        z = tf.matmul(inputs, self.dense.weights[0], transpose_b=True)
        return self.activation(z + self.biases)
```

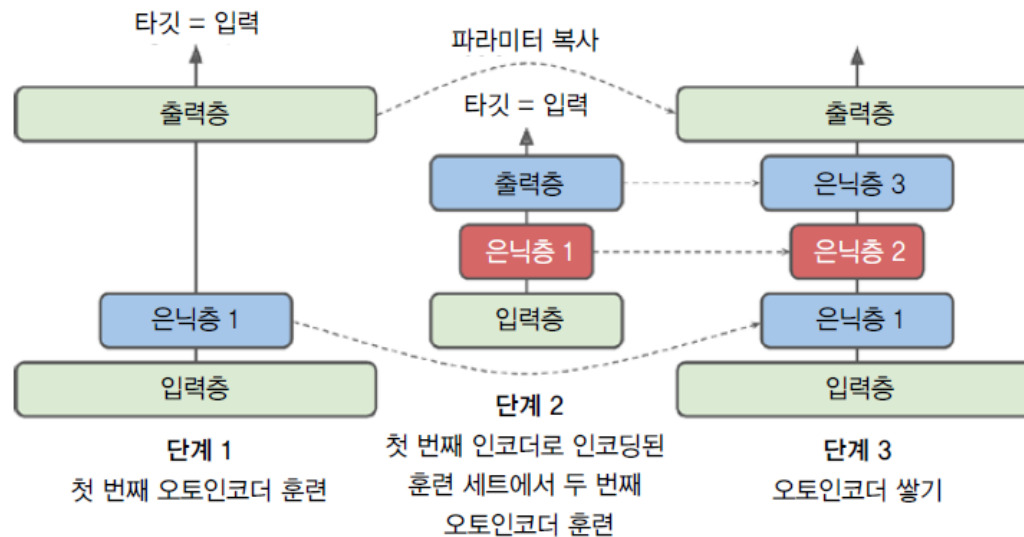
```
dense_1 = keras.layers.Dense(100, activation="selu")
dense_2 = keras.layers.Dense(30, activation="selu")

tied_encoder = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    dense_1,
    dense_2
])

tied_decoder = keras.models.Sequential([
    DenseTranspose(dense_2, activation="selu"),
    DenseTranspose(dense_1, activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])
```

Applications of autoencoder

- Fashion MNIST dataset
 - greedy layerwise training



Convolutional autoencoder

- When you treat the image dataset
 - convolutional network is better than dense network
 - the same idea applies to autoencoders

```
conv_encoder = keras.models.Sequential([
    keras.layers.Reshape([28, 28, 1], input_shape=[28, 28]),
    keras.layers.Conv2D(16, kernel_size=3, padding="SAME", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(32, kernel_size=3, padding="SAME", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2),
    keras.layers.Conv2D(64, kernel_size=3, padding="SAME", activation="selu"),
    keras.layers.MaxPool2D(pool_size=2)
])
conv_decoder = keras.models.Sequential([
    keras.layers.Conv2DTranspose(32, kernel_size=3, strides=2, padding="VALID", activation="selu",
                                  input_shape=[3, 3, 64]),
    keras.layers.Conv2DTranspose(16, kernel_size=3, strides=2, padding="SAME", activation="selu"),
    keras.layers.Conv2DTranspose(1, kernel_size=3, strides=2, padding="SAME", activation="sigmoid"),
    keras.layers.Reshape([28, 28])
])
conv_ae = keras.models.Sequential([conv_encoder, conv_decoder])
```

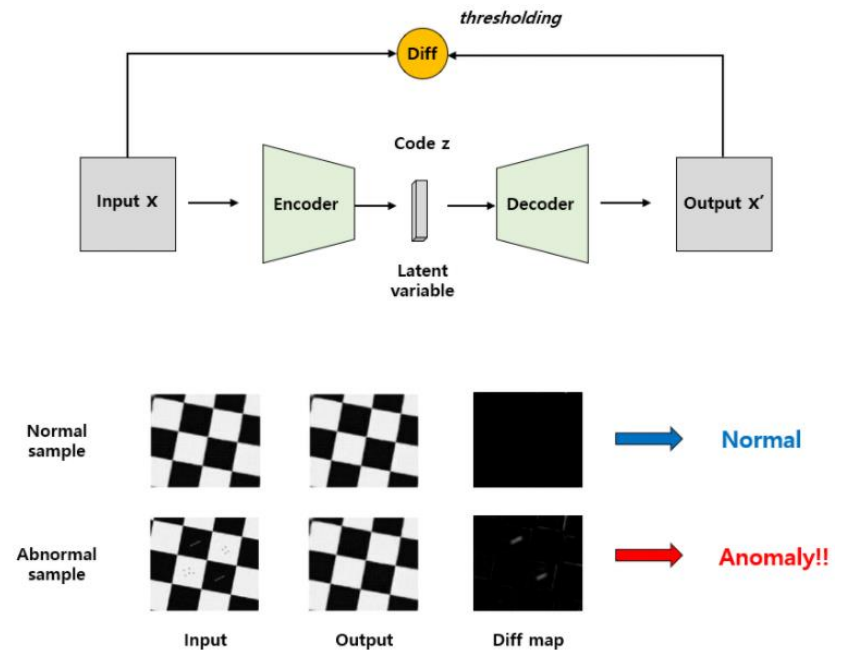

Recurrent autoencoder

- When you treat the time-series dataset
 - recurrent network is better than dense network
 - the same idea applies to autoencoders

```
recurrent_encoder = keras.models.Sequential([
    keras.layers.LSTM(100, return_sequences=True, input_shape=[28, 28]),
    keras.layers.LSTM(30)
])
recurrent_decoder = keras.models.Sequential([
    keras.layers.RepeatVector(28, input_shape=[30]),
    keras.layers.LSTM(100, return_sequences=True),
    keras.layers.TimeDistributed(keras.layers.Dense(28, activation="sigmoid"))
])
recurrent_ae = keras.models.Sequential([recurrent_encoder, recurrent_decoder])
recurrent_ae.compile(loss="binary_crossentropy", optimizer=keras.optimizers.SGD(0.1),
                    metrics=[rounded_accuracy])
```

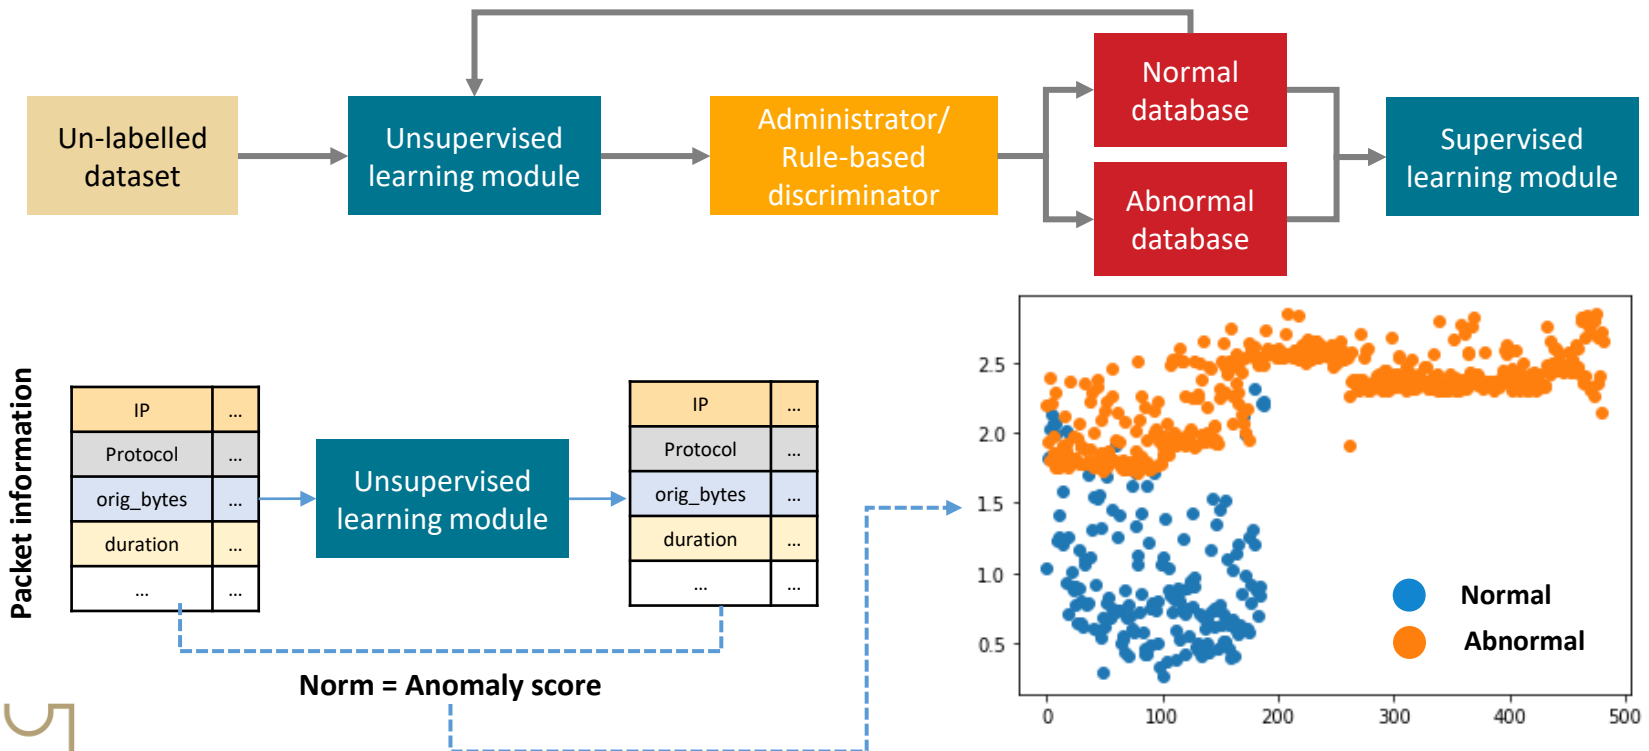
Applications of autoencoder

- Anomaly detection and denoising



Applications of autoencoder

- Anomaly detection and denoising

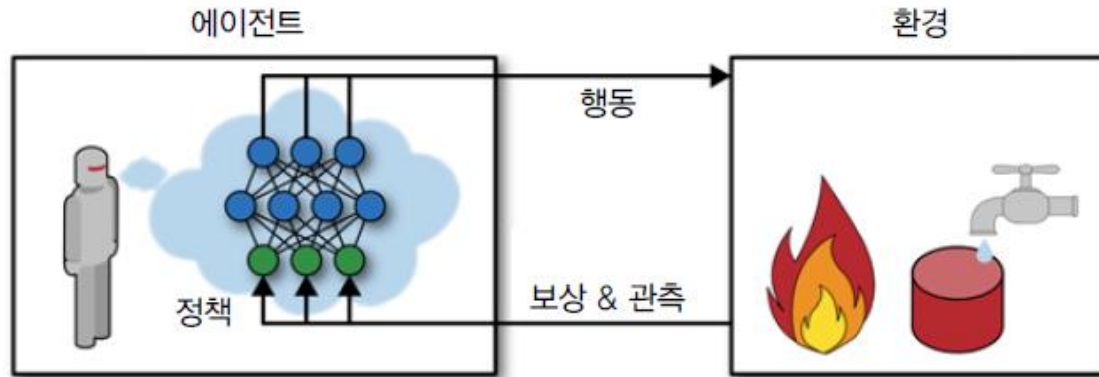


Reinforcement learning

- Definition
 - reinforcement learning is a machine learning training method based on rewarding desired behaviors and/or punishing undesired ones

Reinforcement learning

- Basic structure
 - agent – environment
 - action – reward
 - policy

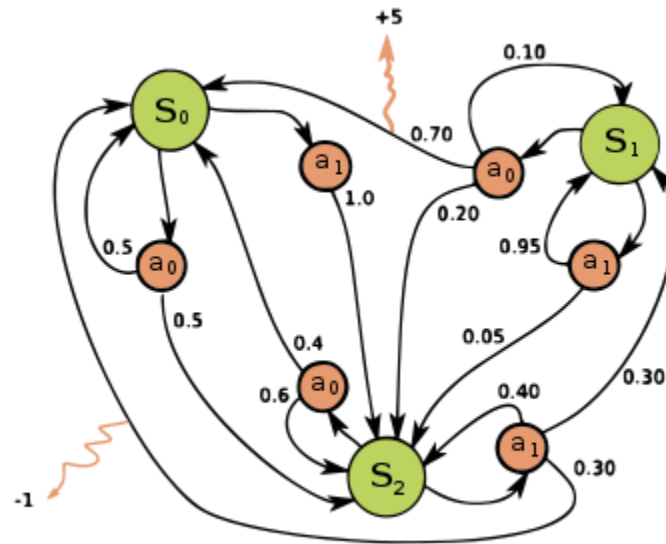


Reinforcement learning

- Sample example
 - https://youtu.be/Yr_nRnqeDp0
 - genetic algorithm is similar with reinforcement learning but they are not identical algorithm


Reinforcement learning

- Markov decision process (MDP)
 - provides a mathematical framework for modeling decision making in situations where outcomes are partly random and partly under the control of a decision maker




Reinforcement learning

- MDP example
 - car acceleration



Reinforcement learning


- State-reward example
 - maze example



- Rewards: -1 per time-step
- Actions: N, E, S, W
- States: Agent's location

Reinforcement learning


- State-reward example
 - maze example



- Arrows represent policy $\pi(s)$ for each state s

Reinforcement learning


- State-reward example
 - maze example



- Numbers represent value $V^\pi(s)$ of each state s

Reinforcement learning

- Q learning
 - q-learning is a model-free reinforcement learning algorithm to learn the value of an action in a particular state



Reinforcement learning

- Q learning
 - formulation
 - action a_t
 - reward r_t
 - state s_t
 - policy $\pi(a|s) = \Pr[a_t = a | s_t = s]$

Reinforcement learning

- Q learning
 - optimization for
 - time-series reward $r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$
 - taking expectation

$$V^\pi(s) = \mathbb{E}_{a_t, a_{t+i}, s_{t+i}} \left[\sum_{i=0}^{\infty} \gamma^i r_{t+i} | s_t = s \right]$$

- by choosing an action

$$Q^\pi(s, a) = \mathbb{E}_{a_{t+i}, s_{t+i}} \left[\sum_{i=0}^{\infty} \gamma^i r_{t+i} | s_t = s, a_t = a \right]$$

Reinforcement learning

- Q learning
 - how to optimize?

$$Q^*(s, a) = \mathbb{E} [r_{t+1} | s_t = s, a_t = a] + \gamma \mathbb{E}_{s_{t+1}} \left[\max_{a'} Q^*(s_{t+1}, a') | s_t = s, a_t = a \right]$$

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
Repeat (for each episode):
 Initialize S
 Repeat (for each step of episode):
 Choose A from S using policy derived from Q (e.g., ϵ -greedy)
 Take action A , observe R, S'
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 $S \leftarrow S'$;
 until S is terminal

Reinforcement learning

- Q learning
 - frozen lake example

S	F	F	F
F	H	F	H
F	F	F	H
H	F	F	G

Reinforcement learning

- Q learning
 - atari game




Reinforcement learning

- Deep Q learning
 - expect Q value through neural network
 1. take some action \mathbf{a}_i and observe $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$, add it to \mathcal{B}
 2. sample mini-batch $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$ from \mathcal{B} uniformly
 3. compute $y_j = r_j + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$ using *target* network $Q_{\phi'}$
 4. $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_\phi(\mathbf{s}_j, \mathbf{a}_j) - y_j)$
 5. update ϕ' : copy ϕ every N steps

Reinforcement learning

- Deep Q learning
 - expect Q value through neural network



Feel free to question
Through e-mail & LMS

본 자료의 연습문제는 수업의 본교재인
한빛미디어, Hands on Machine Learning(2판)에서 주로 발췌함